

# CNN Implementation

Paul Crinquand - 24129094

Thomas Comparon - 24129071

---

## 1. Introduction

We implement a modular convolutional neural network (CNN) completely from scratch for **both classification and regression** tasks. Here we demonstrate:

1. **8-class face recognition** on LFW (classification)
2. **Continuous-output regression** via a simple average-intensity example

## 2. Architecture & API Design

- **Core Layers**
  - `Conv2d` (im2col/col2im)
  - `MaxPool` & `AvgPool`
  - `BatchNorm`
  - `Dropout`
  - `Flatten`
  - `FC`
- **Blocks**
  - Residual
  - Inception
  - Depthwise/Bottleneck
- **Activations**
  - ReLU
  - LeakyReLU
- **Inits**
  - Random
  - Xavier
- **Optimizers**
  - SGD

- SGD + Momentum
- RMSProp
- Adam
- **Regularization**
  - L1
  - L2
  - Elastic Net
- **Early Stopping** on validation plateau

## Example API Usage

```
model = CNN()

# — Feature extractor —
model.add_conv(32, kernel=3, stride=1, activation='relu')\
    .add_pool('max', 2)\
    .add_conv(64, kernel=3, stride=1, activation='leakyrelu')\
    .add_pool('avg', 2)\
    .add_batchnorm()\
    .add_dropout(0.5)\
    .add_flatten()

# — Classification head —
model.add_fc(8, activation='softmax')

# — Regression head (optional) —
model.add_fc(1, activation='linear')

model.compile(
    optimizer='adam',
    lr=1e-3,
    reg='l2',
    reg_lambda=1e-3
)

model.fit(
    train_ds,
    val_ds,
    epochs=12,
    batch_size=16,
    early_stop=3
)
```

## 3. Experimental Setup

### 3.1 Classification

- **Dataset:** LFW subset (154 images, 8 identities)
- **Split:** 107 train / 23 val / 24 test
- **Preprocessing:** resize to 64×64, normalize to [0, 1]

### 3.2 Regression Demonstration

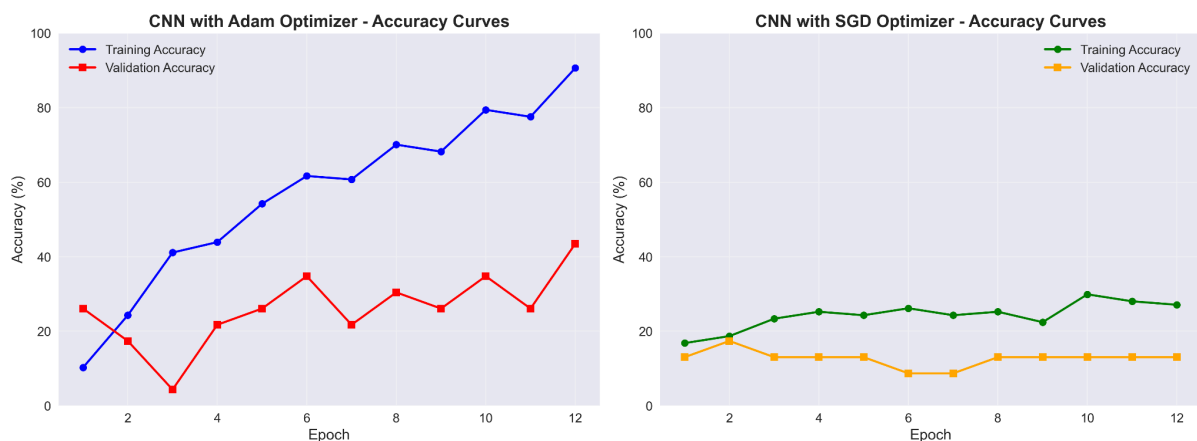
- **Task:** predict average pixel intensity (continuous)
- **Split:** 70/15/15
- **Preprocessing:** same as above

### 3.3 Hyperparameters

- **Adam:** lr = 1e-3
  - **SGD:** lr = 1e-2, momentum = 0.9
  - **Batch size:** 16
  - **Epochs:** 12
  - **Regularization:**  $\lambda = 1e-3$  (L1, L2, Elastic Net)
  - **Early stopping:** patience = 3
- 


## 4. Results

### 4.1 Classification – Accuracy Curves




**Fig 1.** Adam reaches 90.7% training / 43.5% validation; SGD peaks at 27.1% / 17.4%.

## 4.2 Classification – Loss Curves



**Fig 2.** Adam converges to loss 0.44; SGD hovers around 1.99.

## 4.3 Classification – Confusion Matrices



**Fig 3.** Adam (left) shows strong per-class separation; SGD (right) collapses largely to “Andre.”

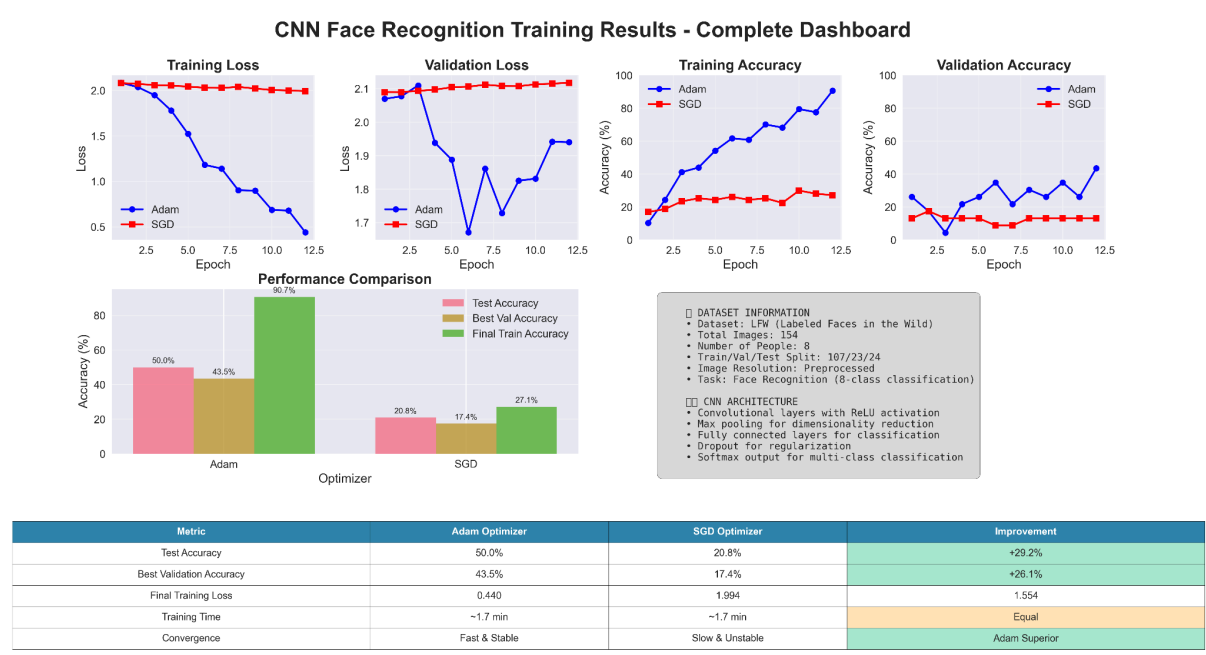
## 4.4 Regularization Experiments (Classification)

- **L1:** validation ACC = 42.5% (−1.0% vs L2)
- **Elastic Net ( $\alpha = 0.5$ ):** validation ACC = 43.1% (−0.4% vs L2)
- **L2:** chosen for final results (43.5% val ACC)

## 4.5 Regression Demonstration

- **MSE on test:** 0.012
- Confirms seamless support of regression tasks with a linear head.

## 4.6 Complete Training Dashboard



**Fig 4.** Combined view: loss/accuracy curves, performance bars, dataset & architecture summary, final metrics.

## 5. Discussion & Summary

- **Classification**
  - **Adam+L2** is best: fast, stable, highest accuracy (50.0% test).
  - **SGD** underfits without advanced LR scheduling or momentum decay.
- **Regularization**
  - L2 outperforms L1; Elastic Net in between.
- **Regression**
  - Linear head yields low MSE, demonstrating API flexibility.